

Monte Carlo Control Algorithm

Name: Meetha Prabhu

Register Number: 212222240065

AIM

To implement Monte Carlo prediction to evaluate an optimal policy in a grid-based environment using Gym's SlippyWalkFive-v0.

PROBLEM STATEMENT

The task involves evaluating the effectiveness of a policy in a grid-based environment using Monte Carlo methods. The environment consists of states and actions, where the goal is to navigate the agent to a terminal state while maximizing rewards. The system needs to determine the action-value and state-value functions for the policy and analyze the policy's performance in terms of success probability and average return.

Monte Carlo Control Algorithm

Step 1:

Initialize Parameters Set up the environment, policy, and initialize the action-value (Q) and state-value (V) functions.

Step 2:

Generate Episodes: Simulate episodes by starting at random states and following the policy until reaching a terminal state.

Step 3:

Compute Returns: Calculate cumulative returns for each state-action pair from the rewards in the episode.

Step 4:

Update Action-Value Function: Average the returns for each state-action pair over

Update Action-Value Function: Average the returns for each state-action pair over multiple episodes to update the action-value function.

Step 5:

Estimate State-Value Function: Derive the state-value function from the action-value function by selecting the best action for each state.

Step 6:

Policy Evaluation: Calculate the success rate and mean return of the policy based on the computed values.

Step 7:

Output: Display the action-value, state-value functions, and the performance metrics for the policy.

MONTE CARLO CONTROL FUNCTION

Import the necessary packages:

```
import warnings ; warnings.filterwarnings('ignore')

import gym, gym_walk
import numpy as np

import random
import warnings

warnings.filterwarnings('ignore', category=DeprecationWarning)
np.set_printoptions(suppress=True)
random.seed(123); np.random.seed(123)
```

Define the printing functions:

```
print(title)
arrs = {k:v for k,v in enumerate(action_symbols)}
for s in range(len(P)):
    a = pi(s)
    print("| ", end="")
    if np.all([done for action in P[s].values() for _, _, done in
action]):
        print("".rjust(9), end=" ")
    else:
        print(str(s).zfill(2), arrs[a].rjust(6), end=" ")
```

```

        if (s + 1) % n_cols == 0: print("|")

def print_state_value_function(V, P, n_cols=4, prec=3, title='State-value
function:'):
    print(title)
    for s in range(len(P)):
        v = V[s]
        print("| ", end="")
        if np.all([done for action in P[s].values() for _, _, done in
action]):
            print("".rjust(9), end=" ")
        else:
            print(str(s).zfill(2), '{}'.format(np.round(v, prec)).rjust(6),
end=" ")
        if (s + 1) % n_cols == 0: print("|")

def probability_success(env, pi, goal_state, n_episodes=100, max_steps=200):
    random.seed(123); np.random.seed(123) ; env.seed(123)
    results = []
    for _ in range(n_episodes):
        state, done, steps = env.reset(), False, 0
        while not done and steps < max_steps:
            state, _, done, h = env.step(pi(state))
            steps += 1
        results.append(state == goal_state)
    return np.sum(results)/len(results)

def mean_return(env, pi, n_episodes=100, max_steps=200):
    random.seed(123); np.random.seed(123) ; env.seed(123)
    results = []
    for _ in range(n_episodes):
        state, done, steps = env.reset(), False, 0
        results.append(0.0)
        while not done and steps < max_steps:
            state, reward, done, _ = env.step(pi(state))
            results[-1] += reward
            steps += 1
    return np.mean(results)

```

Define the environment

```

env = gym.make('FrozenLake-v1')
P = env.env.P
init_state = env.reset()
goal_state = 15
LEFT, RIGHT = range(2)
P

```

Decay Schedule Function:

```

import numpy as np
def decay_schedule(init_value, min_value, decay_ratio, max_steps,
log_start=-2, log_base=10):
    decay_steps = int(max_steps * decay_ratio)
    rem_steps = max_steps - decay_steps
    #This function allows you to calculate all the values for alpha for the full
    training process.
    #(2) First, calculate the number of steps to decay the values using the
    decay_ratio variable. (3) Then, calculate the actual values as an inverse
    log curve. Notice we then normalize between 0 and 1, and finally transform
    the points to lay between init_value and min_value.
    values = np.logspace (log_start, 0, decay_steps, base=log_base,
endpoint=True) [::-1]
    values =(values - values.min()) / (values.max()-values.min())
    values = (init_value - min_value) * values + min_value
    values = np.pad(values, (0, rem_steps), 'edge')
    return values

```

Generate Trajectory:

```

def generate_trajectory(select_action, Q, epsilon, env, max_steps=200): #
Corrected order of arguments
    done, trajectory = False, []
    while not done:
        state = env.reset()
        for t in count():
            action = select_action(state, Q, epsilon)
            next_state, reward, done, _ = env.step(action)
            experience = (state, action, reward, next_state, done)
            trajectory.append(experience)
            if done:
                break
            if t >= max_steps - 1:
                trajectory = []
                break
            state = next_state
    return np.array(trajectory, dtype=object)

```

Monte Carlo Control Function:

```

import numpy as np
from tqdm import tqdm

def mc_control (env, gamma = 1.0,
                init_alpha = 0.5, min_alpha = 0.01, alpha_decay_ratio = 0.5,
                init_epsilon = 1.0, min_epsilon = 0.1, epsilon_decay_ratio =

```

```

        init_epsilon = 1.0, min_epsilon = 0.1, epsilon_decay_ratio =
0.9,

        n_episodes = 3000, max_steps = 200, first_visit = True):
    nS, nA = env.observation_space.n, env.action_space.n

    #Write your code here
    discounts=np.logspace(0,max_steps,num=max_steps, base=gamma,
endpoint=False)
    alphas = decay_schedule(init_alpha, min_alpha,
alpha_decay_ratio,n_episodes)
    epsilons=decay_schedule(init_epsilon, min_epsilon,
epsilon_decay_ratio,n_episodes)
    pi_track=[]
    Q = np.zeros((nS, nA),dtype=np.float64)
    Q_track = np.zeros((n_episodes,nS,nA),dtype=np.float64 )
    select_action = lambda state, Q, epsilon : np.argmax(Q[state]) if
np.random.random()> epsilon else np.random.randint(len(Q[state]))

    for e in tqdm(range(n_episodes),leave=False):
        trajectory = generate_trajectory(select_action,Q, epsilons[e],env,
max_steps)
        visited = np.zeros((nS, nA), dtype=bool)
        for t, (state, action, reward,_,_) in enumerate(trajectory):
            if visited[state][action] and first_visit:
                continue
            visited[state][action]=True
            n_steps=len(trajectory[t:])
            G=np.sum(discounts[:n_steps] * trajectory[t:,2])
            Q[state][action] = Q[state][action] + alphas[e] * (G-Q[state][action])
            Q_track[e]=Q
            pi_track.append(np.argmax(Q,axis=1))
        V=np.max(Q, axis=1)
        pi=lambda s:{s:a for s, a in enumerate(np.argmax(Q, axis=1))} [s]

    # return Q, V, pi, Q_track, pi_track
    return Q, V, pi

```

Print the optimal Value Funtion

```

optimal_Q, optimal_V, optimal_pi = mc_control (env,n_episodes = 3000)
print('Name: Meetha Prabhu      Register Number: 212222240065      ')
print_state_value_function(optimal_Q, P, n_cols=4, prec=2, title='Action-
value function:')
print_state_value_function(optimal_V, P, n_cols=4, prec=2, title='State-
value function:')
print_policy(optimal_pi,P)

```

Probability of Success:

```
# Find the probability of success and the mean return of you your policy
print('Name: Meetha Prabhu      Register Number: 212222240065      ')
print('Reaches goal {:.2f}%. Obtains an average undiscounted return of
{:.4f}.'.format(
    probability_success(env, optimal_pi, goal_state=goal_state)*100,
    mean_return(env, optimal_pi)))
```

OUTPUT:

```
Name: Meetha Prabhu      Register Number: 212222240065
Action-value function:
| 00 [0.11 0.09 0.09 0.07] | 01 [0.02 0.01 0.05 0.08] | 02 [0.05 0.08 0.02 0.01] | 03 [0.06 0.01 0.  0.  ] |
| 04 [0.12 0.05 0.06 0.05] |      | 06 [0.01 0.05 0.11 0.  ] |      |
| 08 [0.06 0.07 0.06 0.13] | 09 [0.05 0.09 0.16 0.09] | 10 [0.2  0.29 0.18 0.02] |      |
|      | 13 [0.01 0.04 0.15 0.22] | 14 [0.14 0.29 0.57 0.3 ] |      |
State-value function:
| 00  0.11 | 01  0.08 | 02  0.08 | 03  0.06 |
| 04  0.12 |      | 06  0.11 |      |
| 08  0.13 | 09  0.16 | 10  0.29 |      |
|      | 13  0.22 | 14  0.57 |      |
Policy:
| 00  < | 01  ^ | 02  v | 03  < |
| 04  < |      | 06  > |      |
| 08  ^ | 09  > | 10  v |      |
|      | 13  ^ | 14  > |      |
```

```
Name: Meetha Prabhu      Register Number: 212222240065
Reaches goal 16.00%. Obtains an average undiscounted return of 0.1600.
```

Mention the Action value function, optimal value function, optimal policy, and success rate for the optimal policy.

RESULT:

Thus the program to implement Monte Carlo control for a given environment is implemented successfully.